

**Course: CS634 - 001**

**Professor: Jason Wang**

**Student Name: Amey Pai Raiker**

**UCID: ap2225**

**Email: ap2225@njit.edu**

# **Wine quality Prediction**

## Table of Contents

Section No	Name	Page no
1	Project Topic Overview	3
2	Dataset	4
3	Data Information	5
4	Data Visualization	6
5	Label Encoder	8
6	Train-Test split	9
7	Model Build	10
8	Confusion matrix and accuracy check	11
9	Project Source code	12
10	Support Vector Machine Source code	16
11	Decision Tree source code	28

## **Project Overview:**

**Problem:** Based on 11 features such as fixed acidity and chlorides etc predict the quality of wine.

## **Data details:**

1. Number of Instances: red wine - 1599

2. Number of Attributes: 11 + output attribute

3. Attribute information:

Input variables (based on physicochemical tests):

- fixed acidity
- volatile acidity
- citric acid
- residual sugar
- chlorides
- free sulfur dioxide
- total sulfur dioxide
- density
- pH
- sulphates
- alcohol
- Output variable (based on sensory data):
  - quality (score between 0 and 10)

4. Missing Attribute Values: None

## Dataset

	A	B	C	D	E	F	G	H	I	J	K	L
1	fixed acid	volatile ac	citric acid	residual su	chlorides	free sulfur	total sulfur	density	pH	sulphates	alcohol	quality
2	7.4	0.7	0	1.9	0.076	11	34	0.9978	3.51	0.56	9.4	5
3	7.8	0.88	0	2.6	0.098	25	67	0.9968	3.2	0.68	9.8	5
4	7.8	0.76	0.04	2.3	0.092	15	54	0.997	3.26	0.65	9.8	5
5	11.2	0.28	0.56	1.9	0.075	17	60	0.998	3.16	0.58	9.8	6
6	7.4	0.7	0	1.9	0.076	11	34	0.9978	3.51	0.56	9.4	5
7	7.4	0.66	0	1.8	0.075	13	40	0.9978	3.51	0.56	9.4	5
8	7.9	0.6	0.06	1.6	0.069	15	59	0.9964	3.3	0.46	9.4	5
9	7.3	0.65	0	1.2	0.065	15	21	0.9946	3.39	0.47	10	7
10	7.8	0.58	0.02	2	0.073	9	18	0.9968	3.36	0.57	9.5	7
11	7.5	0.5	0.36	6.1	0.071	17	102	0.9978	3.35	0.8	10.5	5
12	6.7	0.58	0.08	1.8	0.097	15	65	0.9959	3.28	0.54	9.2	5
13	7.5	0.5	0.36	6.1	0.071	17	102	0.9978	3.35	0.8	10.5	5
14	5.6	0.615	0	1.6	0.089	16	59	0.9943	3.58	0.52	9.9	5
15	7.8	0.61	0.29	1.6	0.114	9	29	0.9974	3.26	1.56	9.1	5
16	8.9	0.62	0.18	3.8	0.176	52	145	0.9986	3.16	0.88	9.2	5
17	8.9	0.62	0.19	3.9	0.17	51	148	0.9986	3.17	0.93	9.2	5
18	8.5	0.28	0.56	1.8	0.092	35	103	0.9969	3.3	0.75	10.5	7
19	8.1	0.56	0.28	1.7	0.368	16	56	0.9968	3.11	1.28	9.3	5
20	7.4	0.59	0.08	4.4	0.086	6	29	0.9974	3.38	0.5	9	4
21	7.9	0.32	0.51	1.8	0.341	17	56	0.9969	3.04	1.08	9.2	6

## Data Information

To get a brief idea about the dataset and how to go about it.

### Code and Output:

```
22### Load wine quality data into Pandas
23df_red = pd.read_csv("C:/Users/imame/Desktop/Data Mining Project/winequality-red.csv") # input
24
25df_red.head() ## To get a peek of the dataset
26
27df_red.info() # to understand total count, no of null values, data type
28
```

```
In [8]: runfile('C:/Users/imame/Desktop/Data Mining Project/Winequality.py', wdir='C:/Users/imame/Desktop/Data Mining Project')
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
```


## Data Pre-processing



### Data Visualization


#### Input code:

```
31 #Here we see that fixed acidity does not give any specification to classify the quality.
32 fig = plt.figure(figsize = (10,6))
33 sns.barplot(x = 'quality', y = 'fixed acidity', data = df_red)
34
35 #Here we see that its quite a downing trend in the volatile acidity as we go higher the quality
36 fig = plt.figure(figsize = (10,6))
37 sns.barplot(x = 'quality', y = 'volatile acidity', data = df_red)
38
39 #Composition of citric acid go higher as we go higher in the quality of the wine
40 fig = plt.figure(figsize = (10,6))
41 sns.barplot(x = 'quality', y = 'citric acid', data = df_red)
42
43 fig = plt.figure(figsize = (10,6))
44 sns.barplot(x = 'quality', y = 'residual sugar', data = df_red)
45
46 #Composition of chloride also go down as we go higher in the quality of the wine
47 fig = plt.figure(figsize = (10,6))
48 sns.barplot(x = 'quality', y = 'chlorides', data = df_red)
49
50 fig = plt.figure(figsize = (10,6))
51 sns.barplot(x = 'quality', y = 'free sulfur dioxide', data = df_red)
```

#### Output:







## Data Pre-processing Label Encoder

Label Encoder to divide the output data into 2 types instead of 10 different floating values.

Input: To use label encoder library

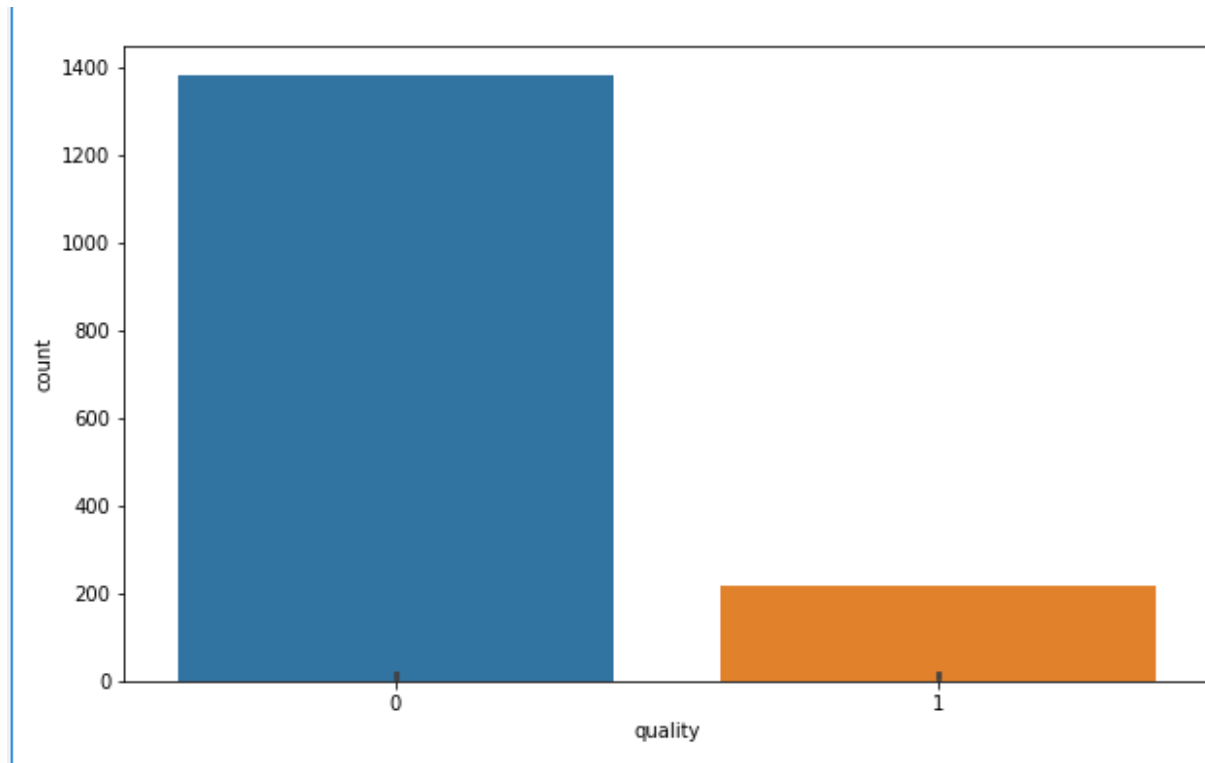
```
6 bins = (2, 6.5, 8)
7 group_names = ['bad', 'good']
8 df_red['quality'] = pd.cut(df_red['quality'], bins = bins, labels = group_names) #
```

To separate values into labels:

```
1 #Now Lets assign a labels to our quality variable
2 label_quality = LabelEncoder()
3
4
5 #Bad becomes 0 and good becomes 1
6 df_red['quality'] = label_quality.fit_transform(df_red['quality']) # To encode the labels as 0 or 1
7 df_red['quality'].value_counts() ## Get the count of each label
8
9 sns.countplot(df_red['quality']) # Craet a bar plot to know the values distribution of the quality
```



**Output:**



Separate the dataset into X and Y variable.

```
1 #Now seperate the dataset as response  
2 X = df_red.drop('quality', axis = 1)  
3 y = df_red['quality'] # get the response variable
```

## Train-Test split

**Splitting the Data into Training and Testing with 80% training data and 20% testing data.**

```
1 #Train and Test splitting of data  
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)  
3 #Applying Standard scaling to get optimized result  
4 sc = StandardScaler() # Normalize the values so that diff between 2 column values isnt significant  
5 X_train = sc.fit_transform(X_train)  
6 X_test = sc.fit_transform(X_test)
```

## Model Build

### Algorithm 1: Decision Trees

I have used the `DecisionTreeClassifier()` to implement the Decision tree algorithm. The source code f which has been attached.

```
85 dtree = DecisionTreeClassifier(criterion = 'entropy', random_state = 0)
86 dtree.fit(X_train, y_train)
87 pred_dtree = dtree.predict(X_test)
88 #Let's see how our model performed
89 print("-----DecisionTreeClassifier-----")
90 print(classification_report(y_test, pred_dtree))
91
92 #Confusion matrix for the random forest classification
93 print(confusion_matrix(y_test, pred_dtree))
94 dtree_cm = confusion_matrix(y_test, pred_dtree)
95 print("Confusion matrix on DecisionTree: ", dtree_cm)
96 dtree_score = accuracy_score(y_test, pred_dtree)
97 print("Accuracy on DecisionTree: ", dtree_score*100)
```

```
DecisionTreeClassifier-----
              precision    recall  f1-score   support

     0       0.91      0.90      0.91       273
     1       0.46      0.49      0.47        47

avg / total       0.84      0.84      0.84       320

[[246  27]
 [ 24  23]]
Confusion matrix on DecisionTree:  [[246  27]
 [ 24  23]]
Accuracy on DecisionTree:  84.0625
```

# Build a Model

## Algorithm 2: Support Vector Machine

I have used the SVC library from sklearn to implement the SVM algorithm for the given dataset. The source code of the decision tree algorithm has been attached.

```
107
108 svc = SVC()
109 svc.fit(X_train, y_train)
110 pred_svc = svc.predict(X_test)
111 print("-----Support Vecotr Machine -----")
112 print(classification_report(y_test, pred_svc))
113 svm_cm = confusion_matrix(y_test, pred_svc)
114 print("Confusion matrix for support vector machine", svm_cm)
115 svm_score = accuracy_score(y_test, pred_svc)
116 print("Accuracy for support vector machine", svm_score*100)
117 #Finding best parameters for our SVC model
```

-----Support Vecotr Machine

	precision	recall	f1-score	support
0	0.88	0.98	0.93	273
1	0.71	0.26	0.37	47
avg / total	0.86	0.88	0.85	320

Confusion matrix for support vector machine [[268 5]  
[ 35 12]]

Accuracy for support vector machine 87.5

## Actual Code for the term project

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

#from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.metrics import accuracy_score


#### Load wine quality data into Pandas

df_red = pd.read_csv("C:/Users/iname/Desktop/Data Mining Project/winequality-red.csv") # input the
red wine dataset

df_red.head() ## To get a peek of the dataset

df_red.info() # to understand total count, no of null values, data type

# Let's do some plotting to know how the data columns are distributed in the dataset

#Here we see that fixed acidity does not give any specification to classify the quality.
fig = plt.figure(figsize = (10,6))
sns.barplot(x = 'quality', y = 'fixed acidity', data = df_red)
```

#Here we see that its quite a downing trend in the volatile acidity as we go higher the quality

```
fig = plt.figure(figsize = (10,6))
```

```
sns.barplot(x = 'quality', y = 'volatile acidity', data = df_red)
```

#Composition of citric acid go higher as we go higher in the quality of the wine

```
fig = plt.figure(figsize = (10,6))
```

```
sns.barplot(x = 'quality', y = 'citric acid', data = df_red)
```

```
fig = plt.figure(figsize = (10,6))
```

```
sns.barplot(x = 'quality', y = 'residual sugar', data = df_red)
```

#Composition of chloride also go down as we go higher in the quality of the wine

```
fig = plt.figure(figsize = (10,6))
```

```
sns.barplot(x = 'quality', y = 'chlorides', data = df_red)
```

```
fig = plt.figure(figsize = (10,6))
```

```
sns.barplot(x = 'quality', y = 'free sulfur dioxide', data = df_red)
```

# Data Pre-processing

#Making binary classificaion for the response variable.

#Dividing wine as good and bad by giving the limit for the quality

```
bins = (2, 6.5, 8)
```

```
group_names = ['bad', 'good']
```

```
df_red['quality'] = pd.cut(df_red['quality'], bins = bins, labels = group_names) #
```

""" pd.cut divides the quality 2-6.5 as bad

6.5 - 8 as good """

```
#Now lets assign a labels to our quality variable
```

```
label_quality = LabelEncoder()
```

```
#Bad becomes 0 and good becomes 1
```

```
df_red['quality'] = label_quality.fit_transform(df_red['quality']) # To encode the labels as 0 or 1
```

```
df_red['quality'].value_counts() ## Get the count of each label
```

```
sns.countplot(df_red['quality']) # Craet a bar plot to know the values distribution of the quality
```

```
#Now seperate the dataset as response variable and feature variabes
```

```
X = df_red.drop('quality', axis = 1) # Get all the columns except the last one
```

```
y = df_red['quality']          # get the last column as the label set
```

```
#Train and Test splitting of data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 42)
```

```
#Applying Standard scaling to get optimized result
```

```
sc = StandardScaler() # Normalize the values so that diff between 2 column values isnt significant
```

```
X_train = sc.fit_transform(X_train)
```

```
X_test = sc.fit_transform(X_test)
```

```
dtree = DecisionTreeClassifier(criterion = 'entropy', random_state = 0)
```

```
dtree.fit(X_train, y_train)
```

```
pred_dtree = dtree.predict(X_test)
```

```
#Let's see how our model performed
```

```
print("-----DecisionTreeClassifier-----")
```

```
print(classification_report(y_test, pred_dtree))
```

```
#Confusion matrix for the random forest classification
```

```

print(confusion_matrix(y_test, pred_dtree))
dtree_cm = confusion_matrix(y_test,pred_dtree)
print("Confusion matrix on DecisionTree: ",dtree_cm)
dtree_score = accuracy_score(y_test,pred_dtree)
print("Accuracy on DecisionTree: ",dtree_score*100)
"""

sgd = SGDClassifier(penalty=None)
sgd.fit(X_train, y_train)
pred_sgd = sgd.predict(X_test)

print(classification_report(y_test, pred_sgd))

print(confusion_matrix(y_test, pred_sgd))
"""

svc = SVC()
svc.fit(X_train, y_train)
pred_svc = svc.predict(X_test)
print("-----Support Vecotr Machine -----")
print(classification_report(y_test, pred_svc))
svm_cm = confusion_matrix(y_test,pred_svc)
print("Confusion matrix for support vector machine",svm_cm)
svm_score = accuracy_score(y_test,pred_svc)
print("Accuracy for support vector machine",svm_score*100)

```

# Code to implement SVM

```
import warnings
import numpy as np
import numpy as np
import libsvm
from libc.stdlib import free

cdef extern from *:
    ctypedef struct svm_parameter:
        pass

np.import_array()

#####
# Internal variables
LIBSVM_KERNEL_TYPES = ['linear', 'poly', 'rbf', 'sigmoid', 'precomputed']

#####
# Wrapper functions

def fit(
    np.ndarray[np.float64_t, ndim=2, mode='c'] X,
    np.ndarray[np.float64_t, ndim=1, mode='c'] Y,
    int svm_type=0, kernel='rbf', int degree=3,
    double gamma=0.1, double coef0=0., double tol=1e-3,
    double C=1., double nu=0.5, double epsilon=0.1,
    np.ndarray[np.float64_t, ndim=1, mode='c']
        class_weight=np.empty(0),
    np.ndarray[np.float64_t, ndim=1, mode='c']
        sample_weight=np.empty(0),
    int shrinking=1, int probability=0,
    double cache_size=100.,
    int max_iter=-1,
    int random_seed=0):
    """
    Train the model using libsvm (low-level method)
    Parameters
    -----
```



`X` : array-like, dtype=float64, size=[`n_samples`, `n_features`]  
`Y` : array, dtype=float64, size=[`n_samples`]  
     target vector  
`svm_type` : {0, 1, 2, 3, 4}, optional  
     Type of SVM: C\_SVC, NuSVC, OneClassSVM, EpsilonSVR or NuSVR  
     respectively. 0 by default.  
`kernel` : {'linear', 'rbf', 'poly', 'sigmoid', 'precomputed'}, optional  
     Kernel to use in the model: linear, polynomial, RBF, sigmoid  
     or precomputed. 'rbf' by default.  
`degree` : int32, optional  
     Degree of the polynomial kernel (only relevant if kernel is  
     set to polynomial), 3 by default.  
`gamma` : float64, optional  
     Gamma parameter in rbf, poly and sigmoid kernels. Ignored by other  
     kernels. 0.1 by default.  
`coef0` : float64, optional  
     Independent parameter in poly/sigmoid kernel. 0 by default.  
`tol` : float64, optional  
     Numeric stopping criterion (WRITEME). 1e-3 by default.  
`C` : float64, optional  
     C parameter in C-Support Vector Classification. 1 by default.  
`nu` : float64, optional  
     0.5 by default.  
`epsilon` : double, optional  
     0.1 by default.  
`class_weight` : array, dtype float64, shape (`n_classes`,), optional  
     np.empty(0) by default.  
`sample_weight` : array, dtype float64, shape (`n_samples`,), optional  
     np.empty(0) by default.  
`shrinking` : int, optional  
     1 by default.  
`probability` : int, optional  
     0 by default.  
`cache_size` : float64, optional  
     Cache size for gram matrix columns (in megabytes). 100 by default.  
`max_iter` : int (-1 for no limit), optional.  
     Stop solver after this many iterations regardless of accuracy  
     (XXX Currently there is no API to know whether this kicked in.)  
     -1 by default.  
`random_seed` : int, optional  
     Seed for the random number generator used for probability estimates.  
     0 by default.  
Returns

```

-----
support : array, shape=[n_support]
    index of support vectors
support_vectors : array, shape=[n_support, n_features]
    support vectors (equivalent to X[support]). Will return an
    empty array in the case of precomputed kernel.
n_class_SV : array
    number of support vectors in each class.
sv_coef : array
    coefficients of support vectors in decision function.
intercept : array
    intercept in decision function
probA, probB : array
    probability estimates, empty array for probability=False
"""

cdef svm_parameter param
cdef svm_problem problem
cdef svm_model *model
cdef const char *error_msg
cdef np.npy_intp SV_len
cdef np.npy_intp nr

if len(sample_weight) == 0:
    sample_weight = np.ones(X.shape[0], dtype=np.float64)
else:
    assert sample_weight.shape[0] == X.shape[0], \
        "sample_weight and X have incompatible shapes: " + \
        "sample_weight has %s samples while X has %s" % \
        (sample_weight.shape[0], X.shape[0])

kernel_index = LIBSVM_KERNEL_TYPES.index(kernel)
set_problem(
    &problem, X.data, Y.data, sample_weight.data, X.shape, kernel_index)
if problem.x == NULL:
    raise MemoryError("Seems we've run out of memory")
cdef np.ndarray[np.int32_t, ndim=1, mode='c'] \
    class_weight_label = np.arange(class_weight.shape[0], dtype=np.int32)
set_parameter(
    &param, svm_type, kernel_index, degree, gamma, coef0, nu, cache_size,
    C, tol, epsilon, shrinking, probability, <int> class_weight.shape[0],
    class_weight_label.data, class_weight.data, max_iter, random_seed)

```

```

error_msg = svm_check_parameter(&problem, &param)
if error_msg:
    # for SVR: epsilon is called p in libsvm
    error_repl = error_msg.decode('utf-8').replace("p < 0", "epsilon < 0")
    raise ValueError(error_repl)

# this does the real work
cdef int fit_status = 0
with nogil:
    model = svm_train(&problem, &param, &fit_status)

# from here until the end, we just copy the data returned by
# svm_train
SV_len = get_l(model)
n_class = get_nr(model)

cdef np.ndarray[np.float64_t, ndim=2, mode='c'] sv_coef
sv_coef = np.empty((n_class-1, SV_len), dtype=np.float64)
copy_sv_coef (sv_coef.data, model)

# the intercept is just model.rho but with sign changed
cdef np.ndarray[np.float64_t, ndim=1, mode='c'] intercept
intercept = np.empty(int((n_class*(n_class-1))/2), dtype=np.float64)
copy_intercept (intercept.data, model, intercept.shape)

cdef np.ndarray[np.int32_t, ndim=1, mode='c'] support
support = np.empty (SV_len, dtype=np.int32)
copy_support (support.data, model)

# copy model.SV
cdef np.ndarray[np.float64_t, ndim=2, mode='c'] support_vectors
if kernel_index == 4:
    # precomputed kernel
    support_vectors = np.empty((0, 0), dtype=np.float64)
else:
    support_vectors = np.empty((SV_len, X.shape[1]), dtype=np.float64)
    copy_SV(support_vectors.data, model, support_vectors.shape)

# TODO: do only in classification
cdef np.ndarray[np.int32_t, ndim=1, mode='c'] n_class_SV
n_class_SV = np.empty(n_class, dtype=np.int32)
copy_nSV(n_class_SV.data, model)

```

```

cdef np.ndarray[np.float64_t, ndim=1, mode='c'] probA
cdef np.ndarray[np.float64_t, ndim=1, mode='c'] probB
if probability != 0:
    if svm_type < 2: # SVC and NuSVC
        probA = np.empty(int(n_class*(n_class-1)/2), dtype=np.float64)
        probB = np.empty(int(n_class*(n_class-1)/2), dtype=np.float64)
        copy_probB(probB.data, model, probB.shape)
    else:
        probA = np.empty(1, dtype=np.float64)
        probB = np.empty(0, dtype=np.float64)
        copy_probA(probA.data, model, probA.shape)
else:
    probA = np.empty(0, dtype=np.float64)
    probB = np.empty(0, dtype=np.float64)

svm_free_and_destroy_model(&model)
free(problem.x)

return (support, support_vectors, n_class_SV, sv_coef, intercept,
        probA, probB, fit_status)

cdef void set_predict_params(
    svm_parameter *param, int svm_type, kernel, int degree, double gamma,
    double coef0, double cache_size, int probability, int nr_weight,
    char *weight_label, char *weight) except *:
    """Fill param with prediction time-only parameters."""

    # training-time only parameters
    cdef double C = .0
    cdef double epsilon = .1
    cdef int max_iter = 0
    cdef double nu = .5
    cdef int shrinking = 0
    cdef double tol = .1
    cdef int random_seed = -1

    kernel_index = LIBSVM_KERNEL_TYPES.index(kernel)

    set_parameter(param, svm_type, kernel_index, degree, gamma, coef0, nu,
                  cache_size, C, tol, epsilon, shrinking, probability,
                  nr_weight, weight_label, weight, max_iter, random_seed)

```

```

def predict(np.ndarray[np.float64_t, ndim=2, mode='c'] X,
            np.ndarray[np.int32_t, ndim=1, mode='c'] support,
            np.ndarray[np.float64_t, ndim=2, mode='c'] SV,
            np.ndarray[np.int32_t, ndim=1, mode='c'] nSV,
            np.ndarray[np.float64_t, ndim=2, mode='c'] sv_coef,
            np.ndarray[np.float64_t, ndim=1, mode='c'] intercept,
            np.ndarray[np.float64_t, ndim=1, mode='c'] probA=np.empty(0),
            np.ndarray[np.float64_t, ndim=1, mode='c'] probB=np.empty(0),
            int svm_type=0, kernel='rbf', int degree=3,
            double gamma=0.1, double coef0=0.,
            np.ndarray[np.float64_t, ndim=1, mode='c']
                class_weight=np.empty(0),
            np.ndarray[np.float64_t, ndim=1, mode='c']
                sample_weight=np.empty(0),
            double cache_size=100.):
    """
    Predict target values of X given a model (low-level method)
    Parameters
    -----
    X : array-like, dtype=float, size=[n_samples, n_features]
    svm_type : {0, 1, 2, 3, 4}
        Type of SVM: C SVC, nu SVC, one class, epsilon SVR, nu SVR
    kernel : {'linear', 'rbf', 'poly', 'sigmoid', 'precomputed'}
        Type of kernel.
    degree : int
        Degree of the polynomial kernel.
    gamma : float
        Gamma parameter in rbf, poly and sigmoid kernels. Ignored by other
        kernels. 0.1 by default.
    coef0 : float
        Independent parameter in poly/sigmoid kernel.
    Returns
    -----
    dec_values : array
        predicted values.
    """
    cdef np.ndarray[np.float64_t, ndim=1, mode='c'] dec_values
    cdef svm_parameter param
    cdef svm_model *model
    cdef int rv

```

```

cdef np.ndarray[np.int32_t, ndim=1, mode='c'] \
    class_weight_label = np.arange(class_weight.shape[0], dtype=np.int32)

set_predict_params(&param, svm_type, kernel, degree, gamma, coef0,
                  cache_size, 0, <int>class_weight.shape[0],
                  class_weight_label.data, class_weight.data)
model = set_model(&param, <int> nSV.shape[0], SV.data, SV.shape,
                 support.data, support.shape, sv_coef.strides,
                 sv_coef.data, intercept.data, nSV.data, probA.data,
probB.data)

#TODO: use check_model
try:
    dec_values = np.empty(X.shape[0])
    with nogil:
        rv = copy_predict(X.data, model, X.shape, dec_values.data)
    if rv < 0:
        raise MemoryError("We've run out of memory")
finally:
    free_model(model)

return dec_values

def predict_proba(
    np.ndarray[np.float64_t, ndim=2, mode='c'] X,
    np.ndarray[np.int32_t, ndim=1, mode='c'] support,
    np.ndarray[np.float64_t, ndim=2, mode='c'] SV,
    np.ndarray[np.int32_t, ndim=1, mode='c'] nSV,
    np.ndarray[np.float64_t, ndim=2, mode='c'] sv_coef,
    np.ndarray[np.float64_t, ndim=1, mode='c'] intercept,
    np.ndarray[np.float64_t, ndim=1, mode='c'] probA=np.empty(0),
    np.ndarray[np.float64_t, ndim=1, mode='c'] probB=np.empty(0),
    int svm_type=0, kernel='rbf', int degree=3,
    double gamma=0.1, double coef0=0.,
    np.ndarray[np.float64_t, ndim=1, mode='c']
        class_weight=np.empty(0),
    np.ndarray[np.float64_t, ndim=1, mode='c']
        sample_weight=np.empty(0),
    double cache_size=100.):
    """
    Predict probabilities
    svm_model stores all parameters needed to predict a given value.

```

For speed, all real work is done at the C level in function `copy_predict (libsvm_helper.c)`.  
 We have to reconstruct model and parameters to make sure we stay in sync with the python object.  
 See `sklearn.svm.predict` for a complete list of parameters.

Parameters

-----

`X` : array-like, dtype=float

`kernel` : {'linear', 'rbf', 'poly', 'sigmoid', 'precomputed'}

Returns

-----

`dec_values` : array  
               predicted values.

"""

`cdef np.ndarray[np.float64_t, ndim=2, mode='c'] dec_values`

`cdef svm_parameter param`

`cdef svm_model *model`

`cdef np.ndarray[np.int32_t, ndim=1, mode='c'] \`

`class_weight_label = np.arange(class_weight.shape[0], dtype=np.int32)`

`cdef int rv`

`set_predict_params(&param, svm_type, kernel, degree, gamma, coef0,`  
                     `cache_size, 1, <int>class_weight.shape[0],`  
                     `class_weight_label.data, class_weight.data)`

`model = set_model(&param, <int> nSV.shape[0], SV.data, SV.shape,`  
                   `support.data, support.shape, sv_coef.strides,`  
                   `sv_coef.data, intercept.data, nSV.data,`  
                   `probA.data, probB.data)`

`cdef np.npy_intp n_class = get_nr(model)`

`try:`

`dec_values = np.empty((X.shape[0], n_class), dtype=np.float64)`

`with nogil:`

`rv = copy_predict_proba(X.data, model, X.shape, dec_values.data)`

`if rv < 0:`

`raise MemoryError("We've run out of memory")`

`finally:`

`free_model(model)`

`return dec_values`

`def decision_function(`

```

np.ndarray[np.float64_t, ndim=2, mode='c'] X,
np.ndarray[np.int32_t, ndim=1, mode='c'] support,
np.ndarray[np.float64_t, ndim=2, mode='c'] SV,
np.ndarray[np.int32_t, ndim=1, mode='c'] nSV,
np.ndarray[np.float64_t, ndim=2, mode='c'] sv_coef,
np.ndarray[np.float64_t, ndim=1, mode='c'] intercept,
np.ndarray[np.float64_t, ndim=1, mode='c'] probA=np.empty(0),
np.ndarray[np.float64_t, ndim=1, mode='c'] probB=np.empty(0),
int svm_type=0, kernel='rbf', int degree=3,
double gamma=0.1, double coef0=0.,
np.ndarray[np.float64_t, ndim=1, mode='c']
    class_weight=np.empty(0),
np.ndarray[np.float64_t, ndim=1, mode='c']
    sample_weight=np.empty(0),
double cache_size=100.):
"""
Predict margin (libsvm name for this is predict_values)
We have to reconstruct model and parameters to make sure we stay
in sync with the python object.
"""

cdef np.ndarray[np.float64_t, ndim=2, mode='c'] dec_values
cdef svm_parameter param
cdef svm_model *model
cdef np.npy_intp n_class

cdef np.ndarray[np.int32_t, ndim=1, mode='c'] \
    class_weight_label = np.arange(class_weight.shape[0], dtype=np.int32)

cdef int rv

set_predict_params(&param, svm_type, kernel, degree, gamma, coef0,
                  cache_size, 0, <int>class_weight.shape[0],
                  class_weight_label.data, class_weight.data)

model = set_model(&param, <int> nSV.shape[0], SV.data, SV.shape,
                  support.data, support.shape, sv_coef.strides,
                  sv_coef.data, intercept.data, nSV.data,
                  probA.data, probB.data)

if svm_type > 1:
    n_class = 1
else:
    n_class = get_nr(model)

```



```

n_class = n_class * (n_class - 1) / 2

try:
    dec_values = np.empty((X.shape[0], n_class), dtype=np.float64)
    with nogil:
        rv = copy_predict_values(X.data, model, X.shape, dec_values.data,
n_class)
        if rv < 0:
            raise MemoryError("We've run out of memory")
finally:
    free_model(model)

return dec_values

def cross_validation(
    np.ndarray[np.float64_t, ndim=2, mode='c'] X,
    np.ndarray[np.float64_t, ndim=1, mode='c'] Y,
    int n_fold, svm_type=0, kernel='rbf', int degree=3,
    double gamma=0.1, double coef0=0., double tol=1e-3,
    double C=1., double nu=0.5, double epsilon=0.1,
    np.ndarray[np.float64_t, ndim=1, mode='c']
        class_weight=np.empty(0),
    np.ndarray[np.float64_t, ndim=1, mode='c']
        sample_weight=np.empty(0),
    int shrinking=0, int probability=0, double cache_size=100.,
    int max_iter=-1,
    int random_seed=0):
    """
    Binding of the cross-validation routine (low-level routine)
    Parameters
    -----
    X : array-like, dtype=float, size=[n_samples, n_features]
    Y : array, dtype=float, size=[n_samples]
        target vector
    svm_type : {0, 1, 2, 3, 4}
        Type of SVM: C SVC, nu SVC, one class, epsilon SVR, nu SVR
    kernel : {'linear', 'rbf', 'poly', 'sigmoid', 'precomputed'}
        Kernel to use in the model: linear, polynomial, RBF, sigmoid
        or precomputed.
    degree : int
        Degree of the polynomial kernel (only relevant if kernel is
        set to polynomial)

```

```

gamma : float
    Gamma parameter in rbf, poly and sigmoid kernels. Ignored by other
    kernels. 0.1 by default.
coef0 : float
    Independent parameter in poly/sigmoid kernel.
tol : float
    Stopping criteria.
C : float
    C parameter in C-Support Vector Classification
nu : float
cache_size : float
random_seed : int, optional
    Seed for the random number generator used for probability estimates.
    0 by default.
Returns
-----
target : array, float
"""

cdef svm_parameter param
cdef svm_problem problem
cdef svm_model *model
cdef const char *error_msg
cdef np.npy_intp SV_len
cdef np.npy_intp nr

if len(sample_weight) == 0:
    sample_weight = np.ones(X.shape[0], dtype=np.float64)
else:
    assert sample_weight.shape[0] == X.shape[0], \
        "sample_weight and X have incompatible shapes: " + \
        "sample_weight has %s samples while X has %s" % \
        (sample_weight.shape[0], X.shape[0])

if X.shape[0] < n_fold:
    raise ValueError("Number of samples is less than number of folds")

# set problem
kernel_index = LIBSVM_KERNEL_TYPES.index(kernel)
set_problem(
    &problem, X.data, Y.data, sample_weight.data, X.shape, kernel_index)
if problem.x == NULL:
    raise MemoryError("Seems we've run out of memory")

```

```

cdef np.ndarray[np.int32_t, ndim=1, mode='c'] \
    class_weight_label = np.arange(class_weight.shape[0], dtype=np.int32)

# set parameters
set_parameter(
    &param, svm_type, kernel_index, degree, gamma, coef0, nu, cache_size,
    C, tol, tol, shrinking, probability, <int>
    class_weight.shape[0], class_weight_label.data,
    class_weight.data, max_iter, random_seed)

error_msg = svm_check_parameter(&problem, &param);
if error_msg:
    raise ValueError(error_msg)

cdef np.ndarray[np.float64_t, ndim=1, mode='c'] target
try:
    target = np.empty((X.shape[0]), dtype=np.float64)
    with nogil:
        svm_cross_validation(&problem, &param, n_fold, <double *> target.data)
finally:
    free(problem.x)

return target

def set_verbosity_wrap(int verbosity):
    """
    Control verbosity of libsvm library
    """
    set_verbosity(verbosity)

```

## Source code for Decision Tree

```
import sys

import math
import pandas as pd

class Node(object):
    def __init__(self, attribute, threshold):
        self.attr = attribute
        self.thres = threshold
        self.left = None
        self.right = None
        self.leaf = False
        self.predict = None

# First select the threshold of the attribute to split set of test data on
# The threshold chosen splits the test data such that information gain is
# maximized
def select_threshold(df, attribute, predict_attr):
    # Convert dataframe column to a list and round each value
    values = df[attribute].tolist()
    values = [float(x) for x in values]
    # Remove duplicate values by converting the list to a set, then sort the
    set
    values = set(values)
    values = list(values)
    values.sort()
    max_ig = float("-inf")
    thres_val = 0
    # try all threshold values that are half-way between successive values in
    this sorted list
    for i in range(0, len(values) - 1):
        thres = (values[i] + values[i+1])/2
        ig = info_gain(df, attribute, predict_attr, thres)
        if ig > max_ig:
            max_ig = ig
            thres_val = thres
    # Return the threshold value that maximizes information gained
    return thres_val
```

```

# Calculate info content (entropy) of the test data
def info_entropy(df, predict_attr):
    # Dataframe and number of positive/negatives examples in the data
    p_df = df[df[predict_attr] == 1]
    n_df = df[df[predict_attr] == 0]
    p = float(p_df.shape[0])
    n = float(n_df.shape[0])
    # Calculate entropy
    if p == 0 or n == 0:
        I = 0
    else:
        I = ((-1*p)/(p + n))*math.log(p/(p+n), 2) + ((-1*n)/(p +
n))*math.log(n/(p+n), 2)
    return I

# Calculates the weighted average of the entropy after an attribute test
def remainder(df, df_subsets, predict_attr):
    # number of test data
    num_data = df.shape[0]
    remainder = float(0)
    for df_sub in df_subsets:
        if df_sub.shape[0] > 1:
            remainder +=
float(df_sub.shape[0]/num_data)*info_entropy(df_sub, predict_attr)
    return remainder

# Calculates the information gain from the attribute test based on a given
threshold
# Note: thresholds can change for the same attribute over time
def info_gain(df, attribute, predict_attr, threshold):
    sub_1 = df[df[attribute] < threshold]
    sub_2 = df[df[attribute] > threshold]
    # Determine information content, and subtract remainder of attributes from
it
    ig = info_entropy(df, predict_attr) - remainder(df, [sub_1, sub_2],
predict_attr)
    return ig

# Returns the number of positive and negative data
def num_class(df, predict_attr):
    p_df = df[df[predict_attr] == 1]
    n_df = df[df[predict_attr] == 0]

```

```

        return p_df.shape[0], n_df.shape[0]

# Chooses the attribute and its threshold with the highest info gain
# from the set of attributes
def choose_attr(df, attributes, predict_attr):
    max_info_gain = float("-inf")
    best_attr = None
    threshold = 0
    # Test each attribute (note attributes maybe be chosen more than once)
    for attr in attributes:
        thres = select_threshold(df, attr, predict_attr)
        ig = info_gain(df, attr, predict_attr, thres)
        if ig > max_info_gain:
            max_info_gain = ig
            best_attr = attr
            threshold = thres
    return best_attr, threshold

# Builds the Decision Tree based on training data, attributes to train on,
# and a prediction attribute
def build_tree(df, cols, predict_attr):
    # Get the number of positive and negative examples in the training data
    p, n = num_class(df, predict_attr)
    # If train data has all positive or all negative values
    # then we have reached the end of our tree
    if p == 0 or n == 0:
        # Create a leaf node indicating it's prediction
        leaf = Node(None, None)
        leaf.leaf = True
        if p > n:
            leaf.predict = 1
        else:
            leaf.predict = 0
        return leaf
    else:
        # Determine attribute and its threshold value with the highest
        # information gain
        best_attr, threshold = choose_attr(df, cols, predict_attr)
        # Create internal tree node based on attribute and it's threshold
        tree = Node(best_attr, threshold)
        sub_1 = df[df[best_attr] < threshold]
        sub_2 = df[df[best_attr] > threshold]
        # Recursively build left and right subtree

```

```

        tree.left = build_tree(sub_1, cols, predict_attr)
        tree.right = build_tree(sub_2, cols, predict_attr)
    return tree

# Given a instance of a training data, make a prediction of healthy or colic
# based on the Decision Tree
# Assumes all data has been cleaned (i.e. no NULL data)
def predict(node, row_df):
    # If we are at a leaf node, return the prediction of the leaf node
    if node.leaf:
        return node.predict

    # Traverse left or right subtree based on instance's data
    if row_df[node.attr] <= node.thres:
        return predict(node.left, row_df)
    elif row_df[node.attr] > node.thres:
        return predict(node.right, row_df)

# Given a set of data, make a prediction for each instance using the Decision Tree
def test_predictions(root, df):
    num_data = df.shape[0]
    num_correct = 0
    for index, row in df.iterrows():
        prediction = predict(root, row)
        if prediction == row['Outcome']:
            num_correct += 1
    return round(num_correct/num_data, 2)

# Prints the tree level starting at given level
def print_tree(root, level):
    print(counter*" ", end="")
    if root.leaf:
        print(root.predict)
    else:
        print(root.attr)
    if root.left:
        print_tree(root.left, level + 1)
    if root.right:
        print_tree(root.right, level + 1)

# Cleans the input data, removes 'Diagnosis' column and adds 'Outcome' column
# where 0 means healthy and 1 means colic
def clean(csv_file_name):
    df = pd.read_csv(csv_file_name, header=None)

```

```

        df.columns = ['K', 'Na', 'CL', 'HCO', 'Endotoxin', 'Anioingap', 'PLA2',
        'SDH', 'GLDH', 'TPP', 'Breath rate', 'PCV', 'Pulse rate', 'Fibrinogen', 'Dimer',
        'FibPerDim', 'Diagnosis']
        # Create new column 'Outcome' that assigns healthy horses a value of 0
        (negative case) and
        # horses with colic a value of 1 (positive case), this makes creating our
        decision tree easier
        df['Outcome'] = 0
        df.loc[df['Diagnosis'] == 'colic.', 'Outcome'] = 1
        df.drop(['Diagnosis'], axis=1 )
        cols = df.columns
        df[cols] = df[cols].apply(pd.to_numeric, errors='coerce')
        return df

def main():
    # An example use of 'build_tree' and 'predict'
    df_train = clean('horseTrain.txt')
    attributes = ['K', 'Na', 'CL', 'HCO', 'Endotoxin', 'Anioingap', 'PLA2',
    'SDH', 'GLDH', 'TPP', 'Breath rate', 'PCV', 'Pulse rate', 'Fibrinogen', 'Dimer',
    'FibPerDim']
    root = build_tree(df_train, attributes, 'Outcome')

    print("Accuracy of test data")
    df_test = clean('horseTest.txt')
    print(str(test_predictions(root, df_test)*100.0) + '%')

if __name__ == '__main__':
    main()

```